

Design system

Design Youtube:

- Repeat the question.
- ask lot of questions.
 - break it down to what they really want you to build
 - what are the requirement.
 - think out loud when you approach the problem.
- example of questions:
 - What parts of YouTube do you want me to design?
 - Obviously there's a lot of things to YouTube, like recommendations and editing content and channels and advertisements and managing payments to people.
 - what part do you want be to focus on.
 - upload videos, or playing them.
 - How much video are we talking about, how much traffic are we talking about.
 - requirement terms of latencies or availability. & ask about the budget as well
 - it shows you what type of trade of are we talking about.

Working Backwards:

- Once you get all the requirement, start designing backwards from the customer experience.

For example, YouTube.

Okay, I need to vend massive amounts of videos all around the world at very low latency and at massive scale.

Well, that probably means I need to use a CDN for my most popular videos at least.

Where does the data come from that feeds that CDN,

- who are those customers
 - are they all around the world.
- what are their use cases
- which use cases do you need to concern yourself with.

Defining scaling requirements

- Is it hundreds of users ? Millions or users ?

- you need horizontal calling.
- how often are users coming ? what transaction rate do you need to support ?
- Also define the scale of the data.
 - Hundreds of videos? Millions?
- Youtube example: millions of users, millions of videos
 - You will need every trick in the book for horizontally scaled servers and data storage.
- Some internal tool might not need this level of complexity, however:
 - Always prefer the simplest solution that will work.
 - Vertical scaling still has its place.

Defining latency requirements

- How fast is fast enough ?
 - use CDN usage.
- Youtube example
 - caching video recommendations.
 - caching video metadata, descriptions.

Defining availability requirement

- How much downtime can you tolerate?
- look for single point of failure and try to replicate it - do a backup.

if you've been explicitly told that this is a high availability service. And they might not tell you, right. Sometimes they might expect you to arrive at these answers yourself.

And again, if you just work backwards

from the customer experience that can inform you as to what the optimal customer experience might look like.

So, if they don't tell you what the latency requirements are say, well, you know, as a customer, I expect pages to render pretty quickly.

I would expect, you know, this webpage to render in under one second.

That means I need to have very fast access to this data locally available, probably in a CDN somewhere.

Think Out Loud

Be Honest

- Don't pretend to know stuff you don't know. That won't end well.
- If you're steered into a direction you're unfamiliar with, say so.
- But don't just give up, let the interviewer help you.
- This is an opportunity to learn.

Defending your Design

- don't get defensive - take feedback constructively.



Cold Start Microservice

Cold start is the time it takes for a serverless function to initialize and execute when it is invoked for the first time or after a period of inactivity

- Warm-Up Strategies
 - Scheduled Tasks:
 - after deployment ping health check of the service to ensure warm up.
 - Or every 1 scheduled tasks to ping the health check.
 - using caching solutions
 - prepopulate when deployment is happening.
- Proactive Scaling.
 - use AWS auto Scaling: it can be configured using gantry.
 - monitor the health of your services and automatically adjust the number of instances based on defined policies.
 - This helps mitigate cold starts by maintaining a minimum number of warm instances.
- Optimizing Container Orchestration
 - Use AWS Fargate: Fargate allows for serverless container orchestration, reducing the likelihood of cold starts by abstracting away the underlying infrastructure.

Problem description:

The cold start happens when you do your very first request after prod deployment OR there have been very long down time with no request.

- In addition in finding and allocating resources to your request, there are 4 steps need to be handled in order to return a response.
 - **Code Download:** from Zip file or download from S3.
 - **Start Execution Environment:** be smart when you choose your Programming language. normally when it comes for lambdas Engineers typically use either `nodejs` or `python` which as less start execution time compared to `Java` for example.
 - **Execute Init Code:** everything before the handler function typically `require` . best to minimise the requirment
 -  Bad Code: `import * as Helpers from "./utils/helpers"` → don't import everything.
 -  Good Code: `import { isOddNumber } from "./utils/helpers"` → only import the thing you need.
 - **Execute Handler Code:** Execute the main `handler` function.
- Note: Cold Start also occurs while scaling up.

Decomposing a Function Execution

Code Download

Start Execution Environment

Execute Init Code

Execute Handler Code

1

2

3

4

Full Cold Start

Partial Cold Start

Warm Start

```
1 //1 - Import required libraries
2 var AWSXRay = require('aws-xray-sdk-core')
3 var captureMySQL = require('aws-xray-sdk-mysql')
4 var mysql = captureMySQL(require('mysql2'))
5 var sequelize = require('sequelize')
6
7 //2 - Extract environment variables from runtime
8 const username = process.env.databaseUser
9 const password = process.env.databasePassword
10 const host = process.env.databaseHost
11
12 //Function Entry Point
13 exports.handler = async (event) => {
14     //3 - Create Database Connection
```

9:23 / 13:16

Strategies to Minimize Cold Start

- ◆ Minimize number of library dependencies
- ◆ Only import what you need
- ◆ Raise Memory Configuration
- ◆ Utilize Provisioned Concurrency

Optimising Docker Images: Best Practices

<https://www.squash.io/how-to-improve-docker-container-performance/>

- Use official Base Images and Alpine ones

- example `FROM node:20-alpine` do not use ubuntu image then install node it will be large image size.
- Minimise the number of Layers example
 - Do this `RUN apt-get update && apt-get install -y package1 package2 package3`
 - Not this

```
RUN apt-get update
RUN apt-get install -y package1
RUN apt-get install -y package2
```

- Use `.dockerignore` to Exclude Unnecessary Files
 - for example exclude `node_modules*/` `logs/` use don't need to copy it only the `src` you need.
- Use Specific Tags for Base Images
 - use specific version instead of `latest` so it won't break.
- Optimising Image size
 - Use a minimal base image: Choose a base image that only includes the necessary dependencies for your application example `From node:20-alpine`.
 - Avoid unnecessary packages and dependencies: Only include the packages and dependencies required by your application.
 - Remove unnecessary files: Clean up any unnecessary files and directories in your image.
- Use multi-stage docker build
 - the one with `From node:21` can do `test`, `lint` & `build` → build means turn `TS` to `JS`.
 - the final image that run into a container will be `From node:21-alpine`, which takes the already built project and CMD `['./listen.js']` listen to it.
 - `COPY --from=deps /workdir/dist/appsapi .` `from` the previous image build
 - Note: `alpine` does not have `tsc` compiler to compile TS to JS.
 - you don't need to have the entire `src` at all only the `dist` built code you need and then run it.
- use default docker cache mechanism in user advantage. the order in DockerFile is important
 - the things that are less likely to be changed will be placed on top. will insure to be cached.
 - the things that more likely to be changed will be placed on the bottom. as they will invalidate the cache until the end of the file commands.

Managing Docker Containers: Tips and Tricks

- Use Appropriate Resource Limits
 - Set the necessary amount of CPU and memory resources.
 - This helps prevent one container from monopolising the resources and affecting the performance of other containers.
 - `docker run --cpus=1 --memory=512m my-container`
- Monitor Resource Usage

- Monitoring the resource usage of your Docker containers is essential to identify bottlenecks and optimise performance.
- `docker stats --all`
- Use Docker Volumes for Persistent Data
 - Docker volumes are a great way to manage persistent data for your containers. By using volumes, you can separate data from the container's filesystem, making it easier to manage and backup.
 - `docker run -v myvolume:/data my-container`
- Clean Up Unused Containers and Images
 - unused containers and images can accumulate and consume valuable disk space. It's important to regularly clean up these unused resources to optimize disk usage
- Utilize Docker Compose for Complex Deployments
 - By using Docker Compose, you can easily define the relationships between different containers and manage their configuration. This simplifies the deployment process and ensures consistency across different environments.

Improving Docker Networking: Strategies for Efficiency

- reach docs

Scaling Docker Applications: Techniques for Performance

- Load Balancing
- Horizontal Scaling
 - Horizontal scaling involves adding more instances of Docker containers to handle increased demand. By horizontally scaling your application, you can distribute the workload across multiple containers
- Vertical Scaling
 - increase of cpu and memory of each container
- Caching
 - use Redis cache to minimize the load on db. as your dockers are stateless.
- Monitoring and Optimisation
 - AWS ECS will provide dashboard to monitor performance.

Container Orchestration with Docker

- What is Container Orchestration?
 - Container orchestration is the process of automating the deployment, management, and scaling of containers. It helps in efficiently running and coordinating multiple containers across a cluster of hosts. Orchestration platforms provide features such as service discovery, scaling, load balancing, high availability, and fault tolerance.
 - it provide load balance and auto scaling functionality out of the box.
- Why Use Container Orchestration?

- Container orchestration simplifies the management of complex containerised applications by abstracting away the underlying infrastructure. Here are some key reasons to use container orchestration
 - **Scalability:** auto scanning: Orchestration platforms enable you to easily scale your applications horizontally by adding or removing containers as per the demand.
 - **High Availability:** Orchestration tools monitor the health of containers and automatically restart or replace failed containers, ensuring high availability of your applications.
 - **Load Balancing:** Orchestration platforms distribute incoming traffic across multiple containers, optimising resource utilisation and improving performance.
 - **Service Discovery:** Orchestration tools provide built-in service discovery mechanisms, allowing containers to find and communicate with each other seamlessly.
 - if they are in the same AWS-ECS cluster , gantry env.

Databases

In SQL DB if how to make search faster ?

- we index the column.

Why SQL Indexes?

- Indexes are used by queries to find data from tables quickly.
- Indexes are created on tables and views.
- with out indexes we will result of query scans.
- example `CREATE Index IX_tblEmployee_Salary ON tblEmployee (SALARY ASC).`

What are the types of Indexes in SQL?

- Clustered Index: determines the physical order of data in a table like `id` of the table type `int Primary Key`
 - you can only have one clustered Index in one table.
 - you can drop the clustered index in `id` and do a composite clustered index example (Gender & Salary) index. just like DynamoDB.
 - the clustered Index determines the order which table rows are sorted.
- Non Clustered Index:

Advantage vs disadvantage of Indexes

- Advantage:
 - `SELECT` clause can benefit from indexes.
 - `DELETE` & `UPDATE` can benefit from indexes as well.
 - `SELECT` with `ORDER BY` the index.

- `GROUP BY`
- Disadvantage:
 - Add additional Disk Space: for non Clustered.
 - `DELETE` , `UPDATE` & `INSERT` : can be slow if we have too many indexes, the record as to be deleted from all the index tables.
 - you have to choose the right amount of indexes on your table.
 - non covering query can be slower than covering query: as the non covering query needs to look up in the actual data from the test of fields of our query is for example `SELECT * ..etc`

Kubernetes (EKS):

- Managed Kubernetes Service: you don't need to worry about the master nodes. AWS will take care of it.
- AWS manages Master Nodes.
- Necessary apps pre-installed.
- Scaling and backups.
- you need to focus on your application worker node only.

How to use EKS?

- Setup or preparation steps.
 - Create AWS account.
 - Create a VPC.
 - Create an IAM role with Security Group.
 - AWS user with least permissions.
- Create Cluster Control Plane: the master node (using the IAM role).
 - choose cluster name, k8s version.
 - choose region and VPC for your cluster.
 - set security for your cluster.
- Create Worker Nodes and connect to cluster.
 - these worker nodes will be some EC2 Instances with cpu and ram configurations.
 - Create as a Node Group (Group of Nodes) for autoscaling.
 - Choose cluster it will attach to.
 - Define Security Group, select instance type, resources
 - Autoscaling configuration → define max and min number of Nodes.

AWS - ECS

Types of ECS

- Serverless with Fargate.
- self managed with EC2:

- you require to manage these EC2, update security and update dependencies.

ECS support auto scaling.

- to handle variable volume.

ECS great with services that requires to be live all time

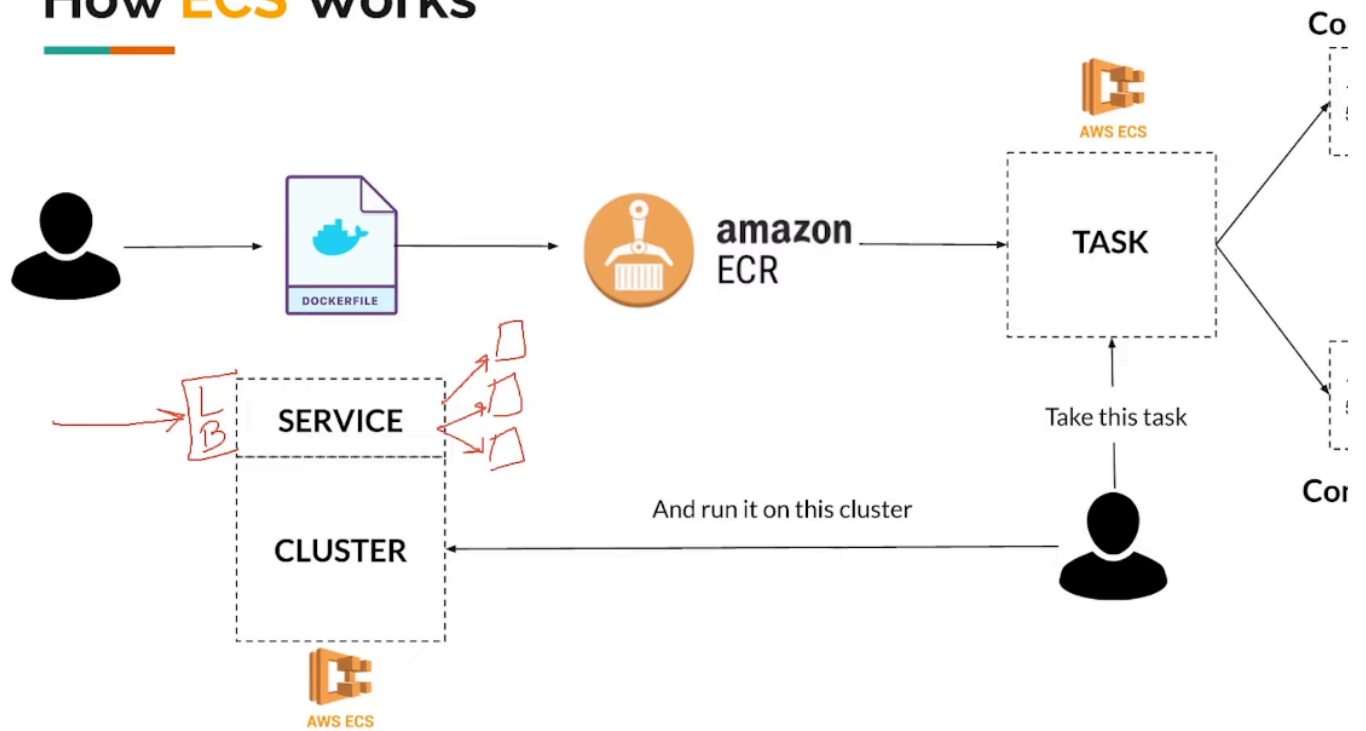
- like a stateless RESTful API.

ECS is cost effective

How ECS works.

- task definition: tells ecs how you want to spin up your docker containers
 - a task can contain more than one container.
 - some applications comes in pair, exmple: rest-api (stateless) & database (stateful).
 - in your task you can specify internal ports than open between these two containers
- Cluster:
 - where your ec2 runners lives: if you use fargate your these ec2 instances will be hidden from you as you don't need to worry about them.
 - you define a **Service** inside your **Cluster**.
 - the **Service** allows you to specify a minimum number of tasks. therefore containers running on this cluster at any point of time.
 - for example: you have very popular application, you need lots of infrastructure to run it on. you would define a service that says at any given point I need at least 10 tasks of this type.
 - what it means you have 10 **Container1** & 10 **Container2**
 - **Service** comes with great tools, it will always monitor your running containers, making sure they are live and healthy.
 - it comes with mentoring dashboard, cpu and memory utilisation.
 - How to make your application scalable?
 - you have application load balancer. and connect to your containers.
 - you have auto scaling solution comes with Service.

How ECS Works



When to move from Monolith to microservice

-